


# Scan images for vulnerabilities

*Estimated reading time: 6 minutes*

✔ These are the docs for DTR version 2.3.11

To select a different version, use the selector below.

2.3.11 ▾



(<https://www.youtube.com/watch?v=121poCB0Nn8>)

Docker Trusted Registry can scan images in your repositories to verify that they are free from known security vulnerabilities or exposures, using Docker Security Scanning. The results of these scans are reported for each image tag.

Docker Security Scanning is available as an add-on to Docker Trusted Registry, and an administrator configures it for your DTR instance. If you do not see security scan results available on your repositories, your organization may not have purchased the Security Scanning feature or it may be disabled. See [Set up Security Scanning in DTR](#)

(<https://docs.docker.com/datacenter/dtr/2.3/guides/admin/configure/set-up-vulnerability-scans/>) for more details.

**Tip:** Only users with write access to a repository can manually start a scan. Users with read-only access can view the scan results, but cannot start a new scan.

# The Docker Security Scan process

Scans run either on demand when a user clicks the **Start a Scan** links or **Scan** button (see Manual scanning (/datacenter/dtr/2.3/guides/user/manage-images/scan-images-for-vulnerabilities/#manual-scanning) below), or automatically on any `docker push` to the repository.

First the scanner performs a binary scan on each layer of the image, identifies the software components in each layer, and indexes the SHA of each component in a bill-of-materials. A binary scan evaluates the components on a bit-by-bit level, so vulnerable components are discovered even if they are statically-linked or under a different name.

The scan then compares the SHA of each component against the US National Vulnerability Database that is installed on your DTR instance. When this database is updated, DTR reviews the indexed components for newly discovered vulnerabilities.

DTR scans both Linux and Windows images, but by default Docker doesn't push foreign image layers for Windows images so DTR can't scan them. If you want DTR to scan your Windows images, configure Docker to always push image layers (https://docs.docker.com/datacenter/dtr/2.3/guides/user/manage-images/pull-and-push-images/), and it will scan the non-foreign layers.

## Security scan on push

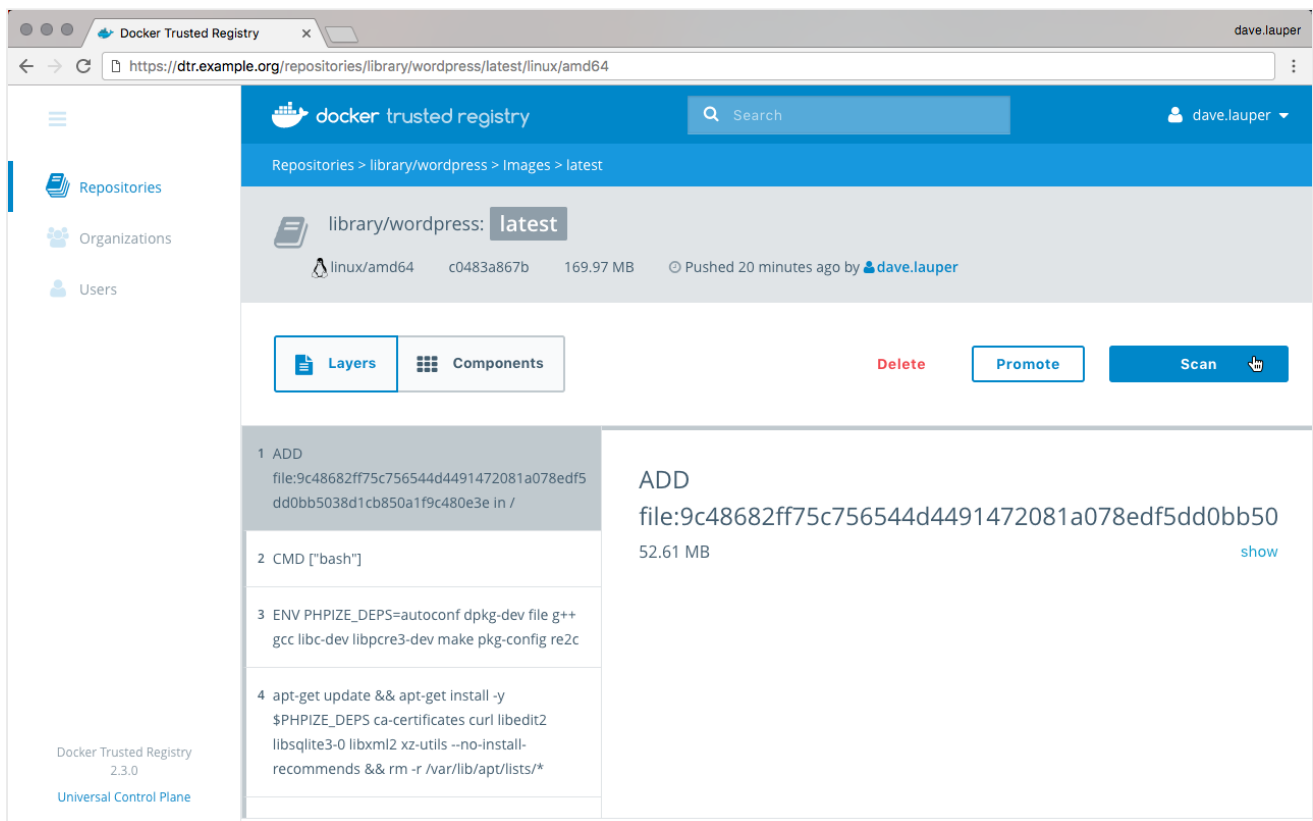
By default, Docker Security Scanning runs automatically on `docker push` to an image repository.

If your DTR instance is configured in this way, you do not need to do anything once your `docker push` completes. The scan runs automatically, and the results are reported in the repository's **Images** tab after the scan finishes.

## Manual scanning

If your repository owner enabled Docker Security Scanning but disabled automatic scanning, you can manually start a scan for images in repositories to which you have `write` access.

To start a security scan, navigate to the **tag details**, and click the **Scan** button.



DTR begins the scanning process. You will need to refresh the page to see the results once the scan is complete.

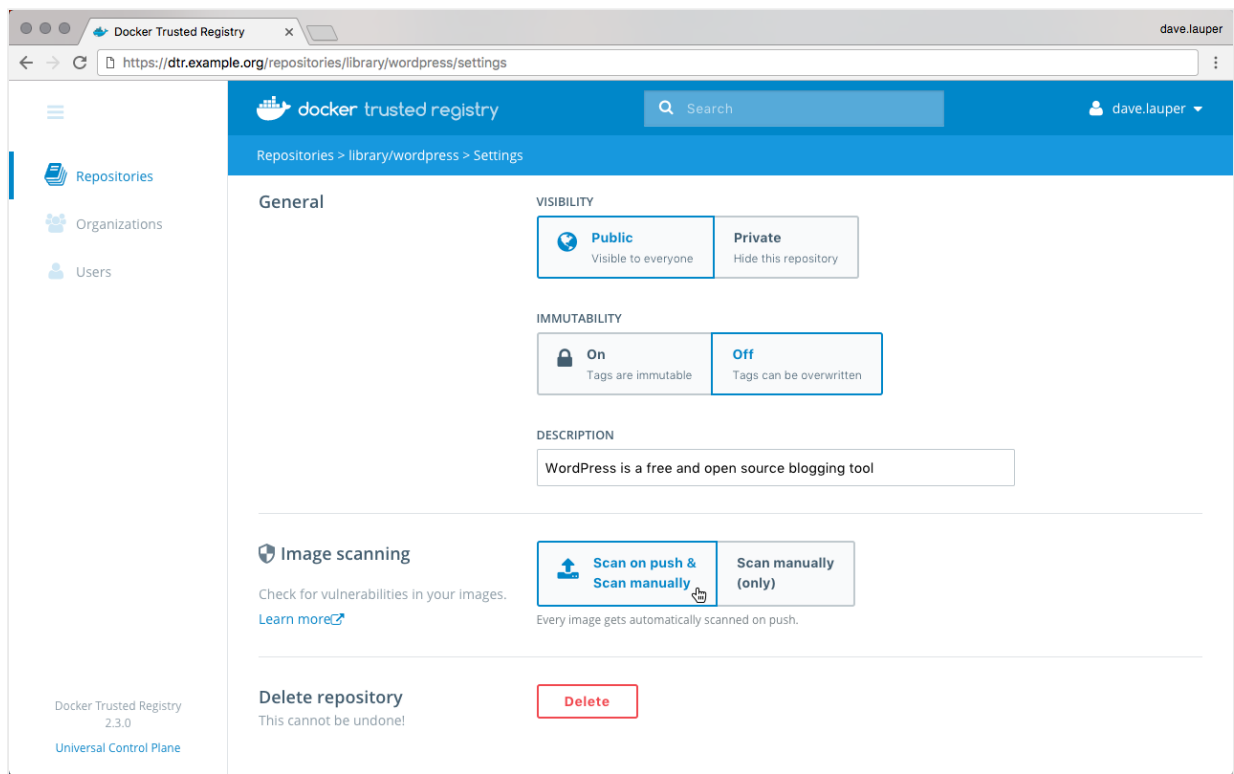
## Change the scanning mode

You can change the scanning mode for each individual repository at any time. You might want to disable scanning if you are pushing an image repeatedly during troubleshooting and don't want to waste resources scanning and re-scanning, or if a repository contains legacy code that is not used or updated frequently.

**Note:** To change an individual repository's scanning mode, you must have `write` or `administrator` access to the repo.

To change the repository scanning mode:

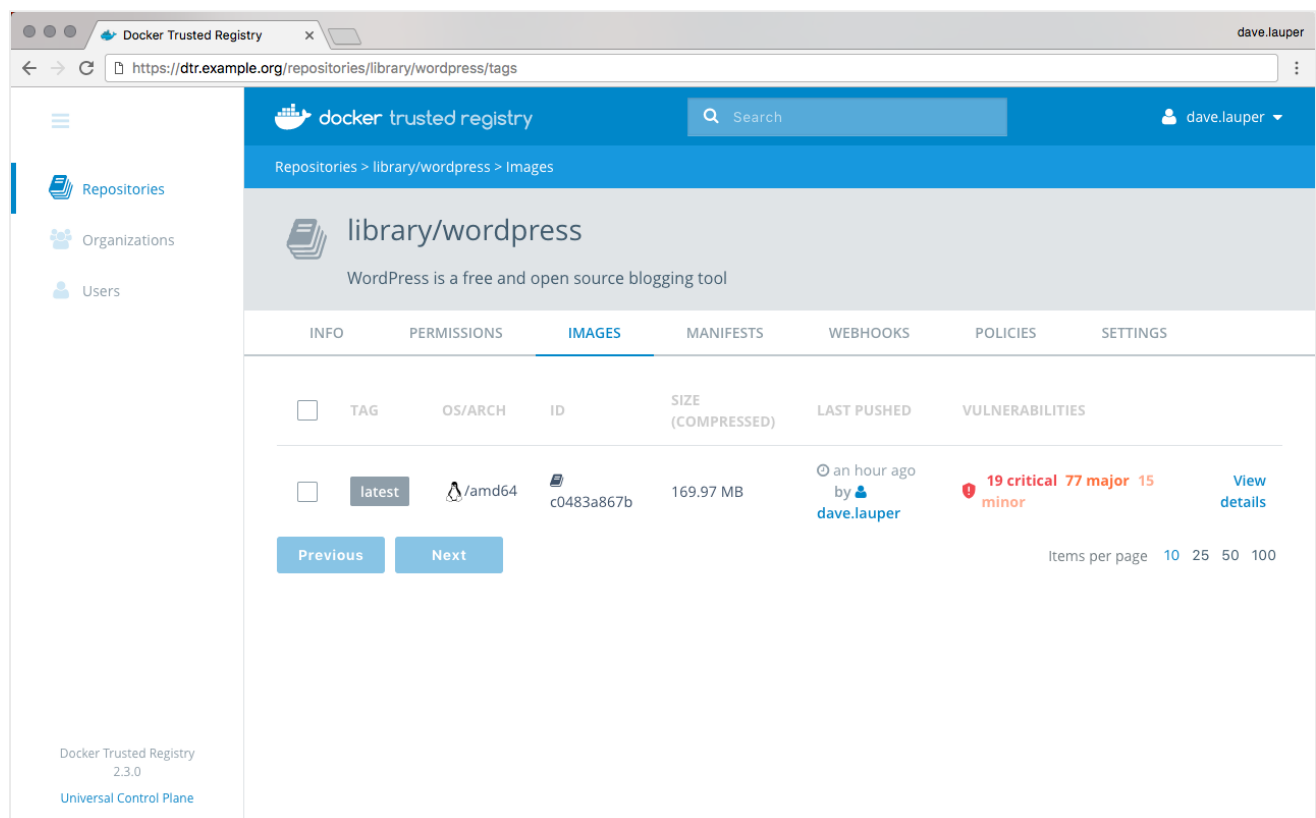
1. Navigate to the repository, and click the **Settings** tab.
2. Scroll down to the **Image scanning** section.
3. Select the desired scanning mode.



## View security scan results

Once DTR has run a security scan for an image, you can view the results.

The **Images** tab for each repository includes a summary of the most recent scan results for each image.



- A green shield icon with a check mark indicates that the scan did not find any vulnerabilities.
- A red or orange shield icon indicates that vulnerabilities were found, and the number of vulnerabilities is included on that same line.

If the vulnerability scan can't detect the version of a component, it reports the vulnerabilities for all versions of that component.

From the **Images** tab you can click **View details** for a specific tag to see the full scan results. The top of the page also includes metadata about the image, including the SHA, image size, date last pushed and user who last pushed, the security scan summary, and the security scan progress.

The scan results for each image include two different modes so you can quickly view details about the image, its components, and any vulnerabilities found.

- The **Layers** view lists the layers of the image in order as they are built by the Dockerfile.

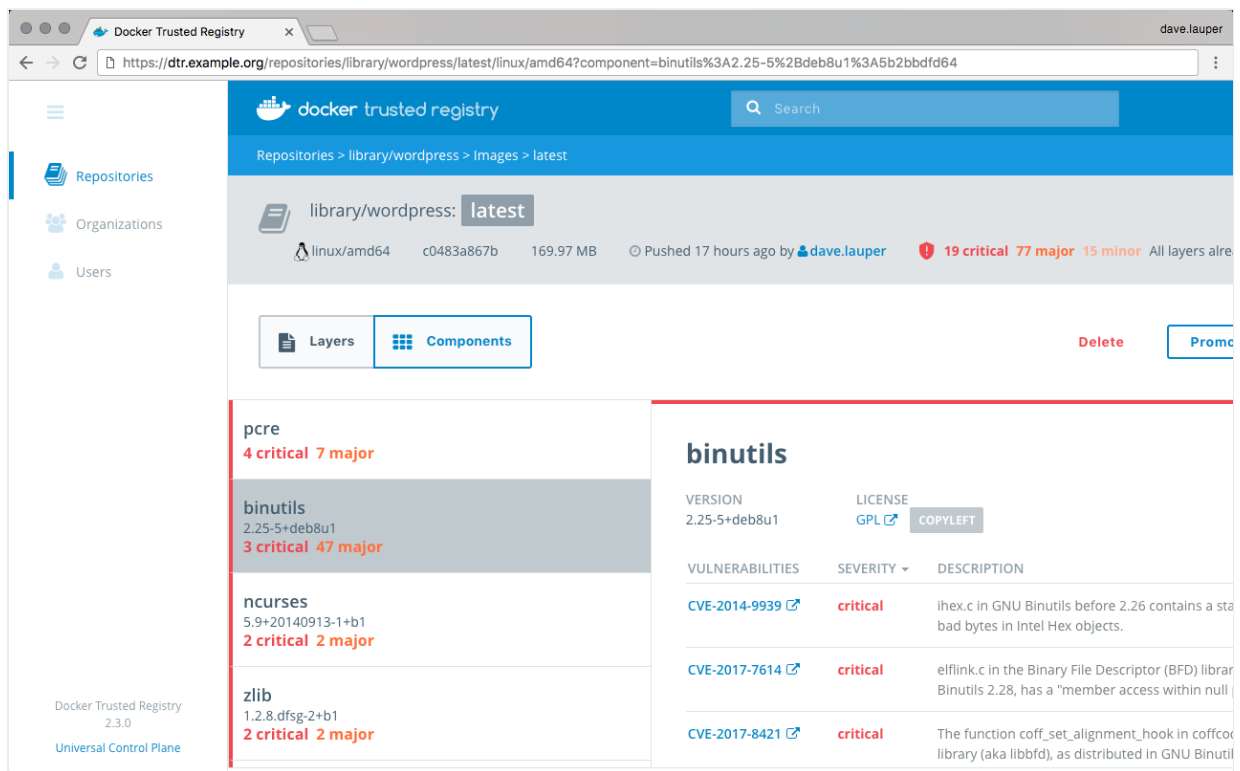
This view can help you find exactly which command in the build introduced the vulnerabilities, and which components are associated with that single command. Click a layer to see a summary of its components. You can then click on a component to switch to the Component view and get more details about the specific item.

**Tip:** The layers view can be long, so be sure to scroll down if you don't immediately see the reported vulnerabilities.

The screenshot shows the Docker Trusted Registry (DTR) interface. The browser address bar displays `https://dtr.example.org/repositories/library/wordpress/latest/linux/amd64`. The interface has a sidebar on the left with 'Repositories', 'Organizations', and 'Users'. The main content area shows the 'library/wordpress:latest' image details. It includes the SHA `c0483a867b`, size `169.97 MB`, and push time 'Pushed an hour ago by dave.lauper'. A summary of vulnerabilities is shown: **19 critical 77 major 15 minor**, with a note 'All layers already scanned'. Below this, the 'Layers' tab is active, displaying a list of build layers. The first layer is highlighted, showing the command 'ADD' and the file path `file:9c48682ff75c756544d4491472081a078edf5dd0bb5038d1cb850a1f9c480e3e in /`. To the right, a summary of components and vulnerabilities is shown, including 'zlib 1.2.8.dfsg-2+b1' with 2 critical and 2 major vulnerabilities, 'ncurses 5.9+20140913-1+b1' with 2 critical and 2 major vulnerabilities, and 'systemd 215-17+deb8u7' with 1 critical and 2 major vulnerabilities.

- The **Components** view lists the individual component libraries indexed by the scanning system, in order of severity and number of vulnerabilities found, most vulnerable first.

Click on an individual component to view details about the vulnerability it introduces, including a short summary and a link to the official CVE database report. A single component can have multiple vulnerabilities, and the scan report provides details on each one. The component details also include the license type used by the component, and the filepath to the component in the image.



## What to do next

If you find that an image in your registry contains vulnerable components, you can use the linked CVE scan information in each scan report to evaluate the vulnerability and decide what to do.

If you discover vulnerable components, you should check if there is an updated version available where the security vulnerability has been addressed. If necessary, you might contact the component's maintainers to ensure that the vulnerability is being addressed in a future version or patch update.

If the vulnerability is in a **base layer** (such as an operating system) you might not be able to correct the issue in the image. In this case, you might switch to a different version of the base layer, or you might find an equivalent, less vulnerable base layer. You might also decide that the vulnerability or exposure is acceptable.

Address vulnerabilities in your repositories by updating the images to use updated and corrected versions of vulnerable components, or by using a different components that provide the same functionality. When you have updated the source code, run a build to create a new image, tag the image, and push the updated image to your DTR instance. You can then re-scan the image to confirm that you have addressed the vulnerabilities.

registry (<https://docs.docker.com/glossary/?term=registry>), scan (<https://docs.docker.com/glossary/?term=scan>), vulnerability (<https://docs.docker.com/glossary/?term=vulnerability>)